

Machine Learning Engineer Nanodegree

Capstone Proposal

Benjamin Rouillé d'Orfeuil

January 2, 2018

Domain Background

Convolutional Neural Networks (CNNs) have successfully been applied in the field of image recognition. These algorithms have proven to be incredibly efficient at classifying images and often outperforms other machine learning algorithms at this task. To illustrate, very accurate predictions can be achieved on the well known MNIST database of handwritten digits¹ and the CIFAR-10 dataset² using very simple network architectures. These datasets, however, are relatively simple. Not only the number of classes is very low (10 digits for the MNIST database and 10 object categories for the CIFAR-10 dataset) but each class is very different from one another.

Image recognition on real world images, on the other hand, requires the building of complex models. Deep CNNs developed for the ImageNet Large Scale Visual Recognition Challenge³ (ILSRVC) is a perfect illustration. Unlike the MNIST and the CIFAR-10 databases, the ILSRVC dataset has a large number of classes and the difference between some of the object categories can be very tenuous. The algorithm needs to be able to discriminate between different breeds of dog or types of snake. It is hence not surprising that the ResNet50 model⁴ developed by Microsoft and that won the 2015 ILSRVC edition has 50 convolutional layers and a total of 168 layers.

It is common practice today to use pre-trained state of the art models to classify photographs. CNNs that have been pre-trained on a large and diverse dataset like ImageNet captures universal features in its early layers that are relevant and useful to most classification problems. The weights of the pre-trained CNNs can then be fine-tuned by continuing training it on the dataset under study.

Problem Statement

Yelp is a social networking site that publishes crowd-sourced reviews about local businesses. About two years ago, Kaggle hosted the Yelp Restaurant Photo Classification challenge (see <https://www.kaggle.com/c/yelp-restaurant-photo-classification>). The problem is that labels are optional when yelp users submit a review leaving some restaurants uncategorized. For this reason, Yelp ask competitors to build a model that automatically predict attribute labels for restaurants using their user-submitted photographs. The goal of this project is to build such a model.

Datasets and Inputs

The photographs and labels for this project can be found in the data section of the competition webpage on Kaggle: <https://www.kaggle.com/c/yelp-restaurant-photo-classification/data>.

¹The MNIST database is available at <http://yann.lecun.com/exdb/mnist/>. A quick analysis of this dataset can be found [here](#).

²The CIFAR-10 dataset can be found at the following url: <https://www.cs.toronto.edu/~kriz/cifar.html>. Predictions on this dataset are presented [here](#).

³The 2017 challenge is described on the ImageNet website: <http://image-net.org/challenges/LSVRC/2017/index> and on Kaggle: <https://www.kaggle.com/c/imagenet-object-localization-challenge>.

⁴The Microsoft team presents their model in the following paper: <https://arxiv.org/pdf/1512.03385.pdf>.

Solution Statement

Benchmark Model

Evaluation Metrics

Project Design