

Machine Learning Engineer Nanodegree

Capstone Project

Benjamin Rouillé d'Orfeuil

May 29, 2018

I. Definition

Project Overview

[Yelp](#) is a social networking site that publishes crowd-sourced reviews about local businesses. About two years ago, Yelp challenged Machine Learning practitioners to build a model that automatically tags restaurants with multiple labels using a dataset of user-submitted photographs. The goal of this project is to develop such a model.

The competition was hosted by [Kaggle](#), a platform where data scientists use their skills to produce the best models for predicting and describing datasets uploaded by companies and users. The Yelp Restaurant Photo Classification competition can be found at this url:

<https://www.kaggle.com/c/yelp-restaurant-photo-classification>

Problem Statement

Back in Spring 2016, when this competition took place, Yelp users were able to upload photographs and write reviews without having to tag the venue with labels. It follows that some restaurants can be left un- or only partially-categorized. There are a lot of photographs uploaded on the Yelp site and, for this reason, business labels need to be predicted by a machine learning algorithm if those have not been selected by users during the submission process.

Each photograph belongs to a business and the task is to predict the business attributes purely from the business photographs. Note that this is a multi-instance multi-label classification problem. Indeed, each business has multiple photographs and predictions need to be done at the business level. Likewise, multiple labels can be assigned to the same business and, as a result, potential dependencies among labels need to be accounted for by the classifier. Classifying real world images is a complex endeavor and, often, deep neural networks are best suited to address such problems.

Metrics

The harmonic mean between precision (p) and recall (r), the F_1 score, is used to evaluate the performance of the algorithm:

$$F_1 = 2 \frac{p \cdot r}{p + r} \text{ where } p = \frac{tp}{tp + fp} \text{ and } r = \frac{tp}{tp + fn}$$

In the above formula, tp , fp and fn denote the true positive, false positive and false negative counts, respectively. In a classification task, $p = 1$ for class i means that every item labeled as belonging to class i does indeed belong to class i whereas $r = 1$ for class i means that every item from class i is labeled as belonging to class i . Though, precision says nothing about the number of items from class i that are mislabeled (fn) and recall says nothing about the number of items that are incorrectly labeled as belonging to class i (fp).

Both precision and recall are obviously relevant for a multi-label classification problem and, for this reason, the F_1 score appears as being the best evaluation metric for this project. A good retrieval algorithm will maximize precision and recall simultaneously and, consequently, good performance on both is favored over extremely good performance on one and poor performance on the other. Note that all Kagglers participating to this competition use the F_1 score as evaluation metric. Thus, models can directly be compared and participants easily ranked.

II. Analysis

Data Exploration

The various datasets and inputs for this competition are available in the data section of the Kaggle competition webpage¹. There are 6 different datasets for this competition:

- `train.csv`;
- `train_photo_to_biz_ids.csv`;
- `test_photo_to_biz_ids.csv`;
- `train_photos.tgz`;
- `test_photos.tgz`;
- and `sample_submission.csv`.

There are 9 different business attributes that are encoded as integer ranging from 0 to 8:

0. `good_for_lunch`;
1. `good_for_dinner`;
2. `takes_reservations`;
3. `outdoor_seating`;
4. `restaurant_is_expensive`;
5. `has_alcohol`;
6. `has_table_service`;
7. `ambience_is_classy`;
8. and `good_for_kids`.

The `train.csv` file provides the list of labels for each business id. As mentioned previously, each business has multiple photographs. The correspondence between the photo id and the business id is given in `train_photo_to_biz_ids.csv` (`test_photo_to_biz_ids.csv`) for the training (test) dataset. The photographs for the training (test) dataset are compressed and combined in `train_photos.tgz` (`test_photos.tgz`). All images have `jpg` format and are named after their photo id. There are 234,842 (237,152) photos and 2,000 (10,000) restaurants in the training (test) dataset². Finally, a template file (`sample_submission.csv`) provides the format for the submission. Predictions enclosed in this file are used by Kaggle to calculate the model performance and in turn rank participants.

Table 1 features a small sample of the training dataset. It relates business id, labels and photos id. The number of photos available for each business is calculated and reported in the last column of Table 1. One can see that this statistic varies from one entry to the other. The histogram of the number of photos

	labels	photos id	# photos
business id			
1000	(1, 2, 3, 4, 5, 6, 7)	[438623, 325966, 227692, 407856, 368729, 16319...	54
1001	(0, 1, 6, 8)	[298536, 20346, 8457, 308694, 349310, 407838, ...	9
100	(1, 2, 4, 5, 6, 7)	[338465, 328433, 243861, 361777, 127198, 46652...	84
1006	(1, 2, 4, 5, 6)	[46472, 341947, 396253, 75316, 42330, 244095, ...	22
1010	(0, 6, 8)	[118251, 219940, 27517, 8578, 148347, 433559, ...	11
101	(1, 2, 3, 4, 5, 6)	[13736, 393696, 286907, 86169, 243460, 254663,...	121
1011	(2, 3, 5, 6)	[372371, 116870, 411981, 208597, 127752, 18839...	70
1012	(1, 2, 3, 5, 6)	[287385, 232258, 388225, 151345, 417121, 32754...	37
1014	(1, 2, 4, 5, 6)	[407910, 33911, 269241, 374218, 256236, 296370...	32
1015	(1, 5, 6, 7)	[456770, 44056, 128542, 373344, 87938, 148452,...	145

Table 1: Training dataset. The labels, photo id and number of photos are given for each business id. This table is extracted from the `eda.ipynb` Jupyter notebook.




Figure 1: Training dataset statistics. Left: histogram of the number of photos per business. Right: bar plot of the label frequency. These figures are extracted from the `eda.ipynb` Jupyter notebook.

per business for the training dataset is shown in the left panel of Figure 1. Some businesses have very few photos (as low as 2 photos) whereas other have thousands of photos. There is on average 117 photographs per business. The label frequency is shown in the right panel of Figure 1. The least used attributes are `good_for_lunch` (#0), `restaurant_is_expensive` (#4) and `ambience_is_classy` (#7) whereas the most used attributes are `has_alcohol` (#5), `has_table_service` (#6) and `good_for_kids` (#8).

Exploratory Visualization

Figure 2 shows 24 randomly chosen photos that has been tagged as `good_for_kids` (#8). One can see that these photos provide rich local business information. Some photos capture the ambience/decor of a place whereas other exhibit the food and drinks that are served. Teaching a computer to understand the context of these photos is clearly not an easy task. Objects can easily be misinterpreted and

¹Visit <https://www.kaggle.com/c/yelp-restaurant-photo-classification/data>.

²The datasets are quite large. Both the training and test tar archive files have a size of about 7 GB.

good_for_kids

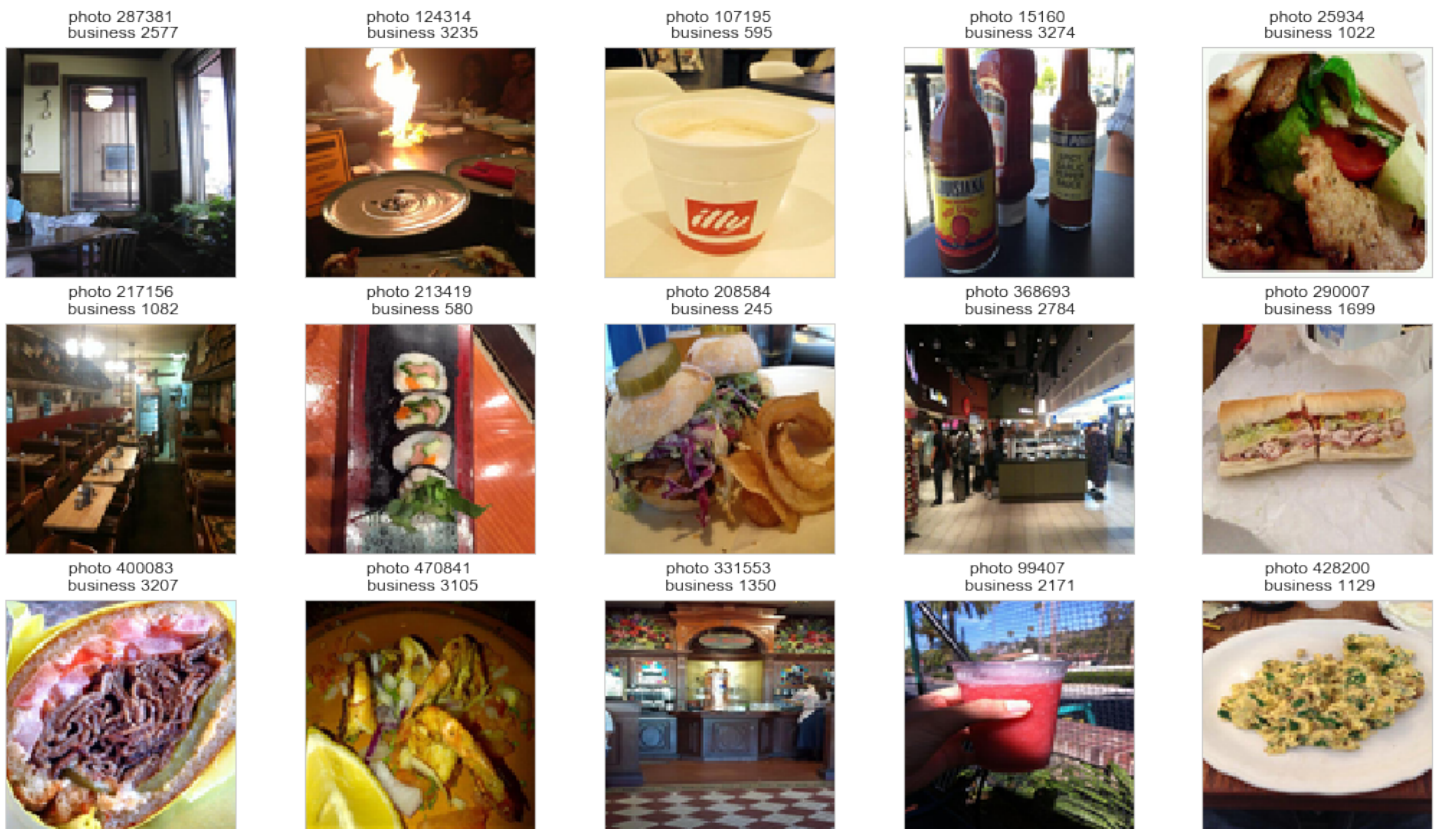


Figure 2: Photos of the training dataset. Each business has been tagged as `good_for_kids` (#8). These images are extracted from the `eda.ipynb` Jupyter notebook.

subsequently misclassified. For example, condiment bottles that can be seen on photo #15160 (1st row and 3rd column) could easily be interpreted as alcohol bottles by a model. Also, it is worth noting that labels are annotated by Yelp users and hence are subjective. To illustrate, some people think that Japanese cuisine is appropriate for kids while other might not. It follows that an establishment that serves sushi (2nd row and 2nd column), for instance, will likely not be labeled as `good_for_kids` (#8) by every Yelp user who has eaten there and submitted photographs.

Algorithms and Techniques

Neural networks have proven to be incredibly efficient at classifying images and often outperform other machine learning algorithms at this task. It comes then as no surprise that deep learning models are used extensively in this project. One now faces two options: i) build and train a deep neural network from scratch or ii) use transfer learning³. The properties of the dataset such as its size and nature usually dictate the type of approach to adopt. The Yelp dataset being both large and complex, it would be unrealistic to train a deep neural network model from scratch given this task would require fine expertise and enormous resources. Also, deep neural networks that have been pre-trained on large and diverse datasets like ImageNet⁴ capture universal features in its early layers that are relevant and useful for most computer vision problems. Thus, leveraging such features allows to reach a better accuracy than any method that would rely only on the available data. For those reasons, transfer learning is a better approach for this project

³Machine learning technique where a model trained on one task is re-purposed on a second related task.

⁴Large visual database designed for use in visual object recognition software research.

The next step entails the selection of the most relevant pre-trained model for the problem domain. Four state-of-the-art deep learning models whose weights have been pre-trained on the ImageNet database are considered here. For each model, the bottleneck features⁵ are extracted and used as inputs of a very simple classifier. Each classifier is independently trained and their performance is evaluated on an unseen set of features. Based on these results, the best pre-trained deep learning model is selected and used as a fixed features extractor.

Before feeding the bottleneck features to a classifier, one needs to address the multi-instance aspect of this project. There are essentially two options: i) derive a feature vector for each instance and combine them accordingly to get one feature vector per restaurant or ii) assign to each instance the label of its corresponding restaurant, proceed to classification and average the output probabilities for each label. Both scenarios are investigated in this project.

Finally, a classifier is trained and predictions are made. Two models are considered for the classification task: i) a multi-layer neural network with a final layer containing one node for each label and ii) XGBoost (gradient boosted decision trees). It is worth noting that a neural network automatically accounts for eventual dependencies among labels because it shares weights for the different label learning tasks. For the other model, label dependencies are handled through classifier chains.

Benchmark

As mentioned above, the evaluation metric for this Kaggle competition is the F_1 score. Participants ranking works as follow: predictions are uploaded to the Kaggle platform, the F_1 score is calculated for each submission and, subsequently, a leaderboard is generated. There are two different leaderboards on Kaggle. A public leaderboard that is calculated using approximately 30% of the test data and a private leaderboard that is based on the other 70%. As long as the competition is ongoing, participants can upload their predictions to Kaggle and consult their standing on the public leaderboard. Once the competition has completed, the private leaderboard reflects the standing. Note that having two leaderboards ensures that participants do not tune their model with the sole purpose of improving their ranking on the public leaderboard. Machine learning is all about building models that generalize well and this property is addressed via the private leaderboard.

As soon as the competition ends both leaderboards are accessible. The performance achieved by other models and reported on the private leaderboard⁶ is used as a benchmark in this project. There are 355 participants in this competition and the 144 best performing models all achieve $F_1 > 75\%$. Note that the best model obtains $F_1 = 83.177\%$.

III. Methodology

Data Preprocessing

Images of the training (test) dataset do not have the same size (height and/or width, scaling) and need to be reshaped/rescaled before being fed to the pre-trained deep learning models. The subtlety is that not all models use the same convention. To illustrate, some models use images with pixel values ranging from 0 to 1 whereas others expect pixel values to be centered but not normalized. That said, there is no need to worry about because [Keras](#), the deep learning framework used for this project, provides a `preprocess_input` function that transform the image to the format required by each model. Apart from that, the data are tidy and require no other preprocessing.

⁵Last feature map before the fully connected layer.

⁶See <https://www.kaggle.com/c/yelp-restaurant-photo-classification/leaderboard>

Before starting the analysis, the training dataset has been split in three dataset: a training, a validation and a test dataset. The models developed in this project are trained on 176,131 images (75% of the original training dataset), validated on 29,355 images (12.5% of the original training dataset) and tested on 29,356 images (12.5% of the original training dataset). These datasets are produced in the `split_data.ipynb` Jupyter notebook.

Implementation

There are two distinct parts in the analysis. The first part consists in selecting the most relevant pre-trained deep learning model for the problem domain. The second part entails developing a model that predicts attribute labels for each business. The various steps that are undertaken to achieve these goals are thoroughly presented below.

Fixed Features Extractor

As discussed earlier, transfer learning is undoubtedly the best approach for this project. Several deep learning models for image classification with weights trained on ImageNet are made available with Keras. Four models are considered here: VGG16⁷, Xception⁸, ResNet50⁹ and InceptionV3¹⁰. All have achieved excellent results in the ILSRVC¹¹. To illustrate, the Xception model achieves 79% (top-1) and 94.5% (top-5) accuracy on the ImageNet validation dataset. It is common practice today to use these models for feature extraction. Only a small fraction of the datasets is used to select the best fixed feature extractor for this competition. The training, validation and test datasets are now made of 20,000, 2,000 and 2,000 images, respectively.

The VGG16, Xception, ResNet50 and InceptionV3 bottleneck features are extracted for the above training, validation and test datasets. This is performed in the `bottleneckFeaturesExtraction.ipynb` Jupyter notebook and explained hereafter. Images in each dataset are first preprocessed and stacked as 4D-arrays or tensors: `(#images,#rows,#columns,#channels)`, where `#images` is the number of images that is fed to the model, `#rows` (`#columns`) is the width (height) of the images in pixel unit and `#channels` is the depth of the images (`RGB = 3`, `grayscale = 1`). As noted previously, the image preprocessing is handled by Keras and depends on each model. The tensors associated to each dataset are then fed to the pre-trained deep learning models and the feature maps located before the last fully connected layer, the bottleneck features, are extracted¹² and saved on disk. Note that the shape of the bottleneck features is specific to each model since the penultimate layer is different from one model to the other. To illustrate, bottleneck features calculated with VGG16 have shape `(7,7,512)` whereas those extracted with ResNet50 have `(1,1,2048)`.

The bottleneck features are fed to a neural network for classification. Its architecture is extremely simple. A pooling layer first spatially reduces the input feature maps which are then flatten using a fully connected layer. The latter contains 9 nodes, one for each label, and is equipped with a `sigmoid` activation function to output probabilities. This classifier is trained on the VGG16, Xception, ResNet50 and InceptionV3 training bottleneck features independently. The loss function is set to `binary_crossentropy` and a stochastic gradient descent (SGD) optimizer with a learning rate of 10^{-4} is used to minimize the weights of the neural networks. At each epoch, when a new set of weights is derived, the loss function is calculated for the corresponding validation bottleneck features. To avoid overfitting, an `EarlyStopping`

⁷See <https://arxiv.org/abs/1409.1556>.

⁸See <https://arxiv.org/abs/1610.02357>.

⁹See <https://arxiv.org/abs/1512.03385>.

¹⁰See <https://arxiv.org/abs/1512.00567>.

¹¹ImageNet Large Scale Visual Recognition Competition: <http://www.image-net.org/challenges/LSVRC>.

¹²The top layer is easily removable with Keras

function is implemented and the training is stopped if the loss function calculated on the validation dataset does not improve after 20 epochs. The weights of the best model is automatically saved on disk. This procedure is carried out in the `pretrainedModelsComparison.ipynb` Jupyter notebook. Figure 3 shows the the value of the loss function (left) and the accuracy metric across epoch for the ResNet50 training (blue) and validation (red) bottleneck features.




Figure 3: Loss function (left) and accuracy across epoch for the ResNet50 training (blue) and validation (red) bottleneck features. These figures are extracted from the `pretrainedModelsComparison.ipynb` Jupyter notebook. Similar plots were produced for the neural networks trained and validated with the VGG16, Xception and InceptionV3.

IV. Results

V. Conclusion